

OpenAI Developers



Builder, Chatting, Connector Registry, and Evaluation loops, for teams that want to ship and iterate faster.

- On the standards side, OpenAI pushed **AGENTS.md (spec)** and participated in the **AAIF (Agentic AI Foundation)** alongside other ecosystem standards like **Model Context Protocol (MCP)** and **Skills**. The value for developers: more portable agent tooling and fewer one-off integrations as the ecosystem converges on shared conventions.

In addition to our work on agents and related standards, we introduced the **Apps SDK**—an open-source framework that extends the Model Context Protocol (MCP) to let developers build UIs alongside their MCP servers, defining both the logic and interactive interface of applications that can run in clients like ChatGPT.

Why it matters: developers can build agents that are less tightly coupled to a single runtime or UI surface, and more easily integrate OpenAI-powered agents into heterogeneous systems.

Open-weight models

In addition to hosted APIs, OpenAI released **open-weight models** designed for transparency, research, and on-prem or self-hosted deployment while retaining strong reasoning and instruction-following capabilities.

- **gpt-oss 120b & 20b** reasoning models designed for self-hosting and on-prem deployments.
- **gpt-oss-safeguard 120b & 20b** safety and policy models intended to run alongside gpt-oss.

Evaluation, tuning, and shipping safely

- **Evals API** for eval-driven development.
- **Reinforcement fine-tuning (RFT)** using programmable graders.
- **Supervised fine-tuning / distillation** for pushing quality down into smaller, cheaper models once you've validated a task with a larger one.
- **Graders** and the **Prompt optimizer** helped teams run a tighter “eval → improve → re-eval” loop.

Wrapping up